

Homework 3

```
[1] import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np

[2] cifar = tf.keras.datasets.cifar10
(x_train, y_train), (x_test, y_test) = cifar.load_data()

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 14s 0us/step

[3] # Convert 2D to 1D array
y_train = y_train.reshape(-1,)
y_test = y_test.reshape(-1,)

[4] # Normalization
x_train, x_test = x_train / 255.0, x_test / 255.0
```

[1] I imported necessary packages.

[2] Download CIFAR10 dataset and assign it to cifar variable. It already has train and test data separated, so assign those to the x and y variables.

[3] Initially, y is 2D array, I do reshape(-1,) to convert it to a 1D array.

[4] Divide the x by 255.0 for a normalization.

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(filters=32, kernel_size=(3,3), padding='same', activation='relu', input_shape=(32, 32, 3)),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv2D(filters=32, kernel_size=(3,3), padding='same', activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPool2D((2, 2)),
    tf.keras.layers.Dropout(0.2),

    tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3), padding='same', activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv2D(filters=64, kernel_size=(3,3), padding='same', activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPool2D((2, 2)),
    tf.keras.layers.Dropout(0.3),

    tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3), padding='same', activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv2D(filters=128, kernel_size=(3,3), padding='same', activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.MaxPool2D((2, 2)),
    tf.keras.layers.Dropout(0.4),

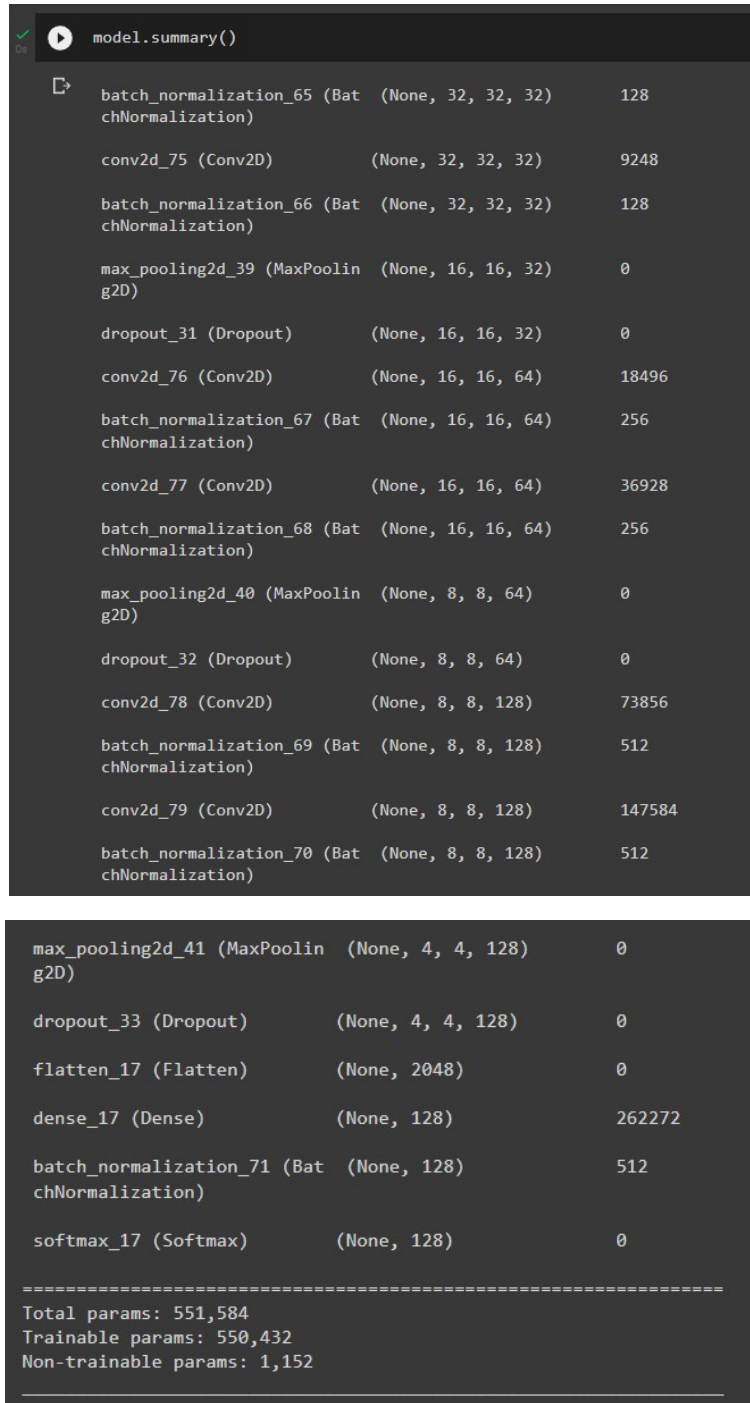
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Softmax()
])

loss = tf.keras.losses.SparseCategoricalCrossentropy()

model.compile(loss=loss, optimizer='adam', metrics=['accuracy'])

tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir='log')
```

This is the Keras CNN model. The input shape is (32, 32, 3). Since I will train in batch, so I have to add batch normalizations. Batch normalization has to be after Conv2D but before MaxPool2D. I train Conv2D → BatchNormalization → Conv2D → Batch Normalization → MaxPool2D → Dropout, and increase the filters size from time to time. I have to deal with the overfitting on training data problem, so I add Dropout layer to my neural network. Then, flatten the layer, add one normal hidden layer, do the batch normalization again, then add the Softmax layer lastly. The loss function is Sparse categorical crossentropy. Optimizer will be AdamOptimizer. Use accuracy metrics for measuring training accuracy.



```
model.summary()
```

batch_normalization_65 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_75 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_66 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d_39 (MaxPooling2D)	(None, 16, 16, 32)	0
dropout_31 (Dropout)	(None, 16, 16, 32)	0
conv2d_76 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_67 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_77 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_68 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_40 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_32 (Dropout)	(None, 8, 8, 64)	0
conv2d_78 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_69 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_79 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_70 (Batch Normalization)	(None, 8, 8, 128)	512
max_pooling2d_41 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_33 (Dropout)	(None, 4, 4, 128)	0
flatten_17 (Flatten)	(None, 2048)	0
dense_17 (Dense)	(None, 128)	262272
batch_normalization_71 (Batch Normalization)	(None, 128)	512
softmax_17 (Softmax)	(None, 128)	0

=====		
Total params:	551,584	
Trainable params:	550,432	
Non-trainable params:	1,152	

These figures show the model summary.

```
[59] model.fit(x_train, y_train, batch_size=64, epochs=100, callbacks=[tensorboard_callback], validation_data=(x_test, y_test))

782/782 [=====] - 8s 10ms/step - loss: 0.1445 - accuracy: 0.9501 - val_loss: 0.4680 - val_accuracy: 0.8662
Epoch 73/100
782/782 [=====] - 8s 10ms/step - loss: 0.1453 - accuracy: 0.9503 - val_loss: 0.4391 - val_accuracy: 0.8755
Epoch 74/100
782/782 [=====] - 8s 10ms/step - loss: 0.1442 - accuracy: 0.9511 - val_loss: 0.4468 - val_accuracy: 0.8711
Epoch 75/100
782/782 [=====] - 8s 10ms/step - loss: 0.1358 - accuracy: 0.9545 - val_loss: 0.4496 - val_accuracy: 0.8707
Epoch 76/100
782/782 [=====] - 8s 10ms/step - loss: 0.1408 - accuracy: 0.9523 - val_loss: 0.4698 - val_accuracy: 0.8676
Epoch 77/100
782/782 [=====] - 8s 10ms/step - loss: 0.1349 - accuracy: 0.9530 - val_loss: 0.4755 - val_accuracy: 0.8672
Epoch 78/100
782/782 [=====] - 8s 10ms/step - loss: 0.1351 - accuracy: 0.9545 - val_loss: 0.4468 - val_accuracy: 0.8750
Epoch 79/100
782/782 [=====] - 8s 10ms/step - loss: 0.1322 - accuracy: 0.9537 - val_loss: 0.4672 - val_accuracy: 0.8729
Epoch 80/100
782/782 [=====] - 8s 10ms/step - loss: 0.1361 - accuracy: 0.9533 - val_loss: 0.4710 - val_accuracy: 0.8725
Epoch 81/100
782/782 [=====] - 8s 10ms/step - loss: 0.1341 - accuracy: 0.9534 - val_loss: 0.4841 - val_accuracy: 0.8656
Epoch 82/100
782/782 [=====] - 8s 10ms/step - loss: 0.1342 - accuracy: 0.9542 - val_loss: 0.4821 - val_accuracy: 0.8674
Epoch 83/100
782/782 [=====] - 8s 10ms/step - loss: 0.1307 - accuracy: 0.9556 - val_loss: 0.4905 - val_accuracy: 0.8660
Epoch 84/100
782/782 [=====] - 8s 10ms/step - loss: 0.1305 - accuracy: 0.9554 - val_loss: 0.4726 - val_accuracy: 0.8683
Epoch 85/100
782/782 [=====] - 8s 10ms/step - loss: 0.1270 - accuracy: 0.9566 - val_loss: 0.4663 - val_accuracy: 0.8714
Epoch 86/100
782/782 [=====] - 8s 10ms/step - loss: 0.1301 - accuracy: 0.9542 - val_loss: 0.4691 - val_accuracy: 0.8709
Epoch 87/100
782/782 [=====] - 8s 10ms/step - loss: 0.1275 - accuracy: 0.9570 - val_loss: 0.4646 - val_accuracy: 0.8714
Epoch 88/100
782/782 [=====] - 8s 10ms/step - loss: 0.1257 - accuracy: 0.9563 - val_loss: 0.4762 - val_accuracy: 0.8698
Epoch 89/100
```

Then, fit the model. My batch size is 64, with 100 epochs. The validation data parameter here uses the variable `x_test` and `y_test` that we assigned it beforehand. Also use Google Colab's GPU because if not, its going to take too long time. After the training is completed, we will check the results.

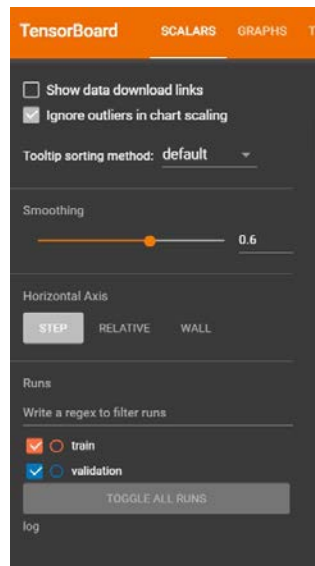
```
model.evaluate(x_train, y_train)

1563/1563 [=====] - 6s 4ms/step - loss: 0.0066 - accuracy: 0.9995
[0.0065909819677472115, 0.9995200037956238]


[61] model.evaluate(x_test, y_test)

313/313 [=====] - 1s 4ms/step - loss: 0.4731 - accuracy: 0.8749
[0.47305503487586975, 0.8748999834060669]
```


Here you can see that the training loss is 0.0066, the training accuracy is 0.9995, the testing loss is 0.4731, and the testing accuracy is 0.8749. I tried changing parameters for a several times, and this one was my best result. Now, let's see the Tensorboard graph results.



Orange line means training, blue line means validation (testing)



Here is the accuracy in each epoch of training (orange) and testing (blue) data.



Here is the loss in each epoch of training (orange) and testing (blue) data.

Here is the visualization of CNN architecture.

